

Complete Database Management Systems (DBMS) & SQL Training Guide

Comprehensive Training Material - Database Fundamentals to Advanced SQL

Table of Contents

1. [Introduction to DBMS & SQL](#)
 2. [Understanding RDBMS](#)
 3. [Database Normalization](#)
 4. [SQL Data Types](#)
 5. [Practical Examples & Best Practices](#)
-

Section 1: Introduction to DBMS & SQL

1.1 What is a Database?

A **database** is an organized collection of structured information or data, typically stored electronically in a computer system. Think of it as a digital filing cabinet where data is organized in a way that makes it easy to access, manage, and update.

Real-World Analogy:

- A library catalog system that tracks books, members, borrowing history
- An e-commerce system storing products, customers, orders
- A hospital system managing patients, doctors, appointments, medical records

1.2 What is a Database Management System (DBMS)?

A **DBMS** is software that interacts with users, applications, and the database itself to capture and analyze data. It provides a systematic way to create, retrieve, update, and manage data.

Key Functions of DBMS:

- **Data Storage:** Efficiently stores large amounts of data
- **Data Retrieval:** Quickly fetches requested data

- **Data Security:** Controls access and protects data
- **Data Integrity:** Ensures accuracy and consistency
- **Concurrent Access:** Multiple users can access data simultaneously
- **Backup & Recovery:** Protects against data loss

Popular DBMS Systems:

- **MySQL** - Open-source, widely used for web applications
- **Microsoft SQL Server** - Enterprise-grade, Windows-integrated
- **Oracle Database** - High-performance, enterprise solutions
- **PostgreSQL** - Advanced open-source with strong standards compliance
- **SQLite** - Lightweight, embedded in applications
- **MongoDB** - NoSQL document database
- **Microsoft Access** - Desktop database for small businesses

1.3 What is SQL?

SQL (Structured Query Language) is the standard programming language specifically designed for managing and manipulating relational databases.

SQL Definition:

SQL = Structured Query Language

Purpose: Retrieval, manipulation, and management of data in relational databases

Nature: Declarative language (you specify WHAT you want, not HOW to get it)

Key Characteristics:

- **Standard Language:** Works across different database systems
- **Declarative:** Focus on results, not implementation
- **Set-Based:** Operates on sets of data, not individual records
- **Case-Insensitive:** `SELECT`, `select`, and `Select` are equivalent (in most systems)
- **English-Like Syntax:** Easy to read and understand

1.4 Why Learn SQL?

1. Universal Standard

- Every major database system uses SQL
- Skills transfer across different platforms
- Industry-standard for 50+ years

2. High Demand

- Required skill for: Data Analysts, Database Administrators, Backend Developers, Data Scientists, Business Intelligence Professionals
- One of the most sought-after technical skills in job market

3. Powerful Data Analysis

- Extract insights from millions of records
- Complex queries in simple syntax
- Generate reports and analytics

4. Database Independence

- Core SQL concepts work across all RDBMS
- Easy to switch between MySQL, SQL Server, Oracle, etc.

5. Career Opportunities

- Database Administrator (DBA)
- Data Analyst
- Backend Developer
- Data Engineer
- Business Intelligence Analyst
- Data Scientist

1.5 SQL Dialects

While standard SQL (ANSI SQL) is the foundation, different database systems have their own extensions and variations:

Major SQL Dialects:

Database System	SQL Dialect	Key Features
Microsoft SQL Server	T-SQL (Transact-SQL)	Stored procedures, variables, error handling, transactions
Oracle	PL/SQL (Procedural Language/SQl)	Packages, cursors, triggers, advanced procedural features
MySQL	MySQL SQL	Auto-increment, LIMIT clause, specific functions
PostgreSQL	PostgreSQL SQL	Advanced features, JSON support, window functions
Microsoft Access	JET SQL	Desktop-oriented, simplified syntax
SQLite	SQLite SQL	Lightweight, embedded, simplified data types

Example of Dialect Differences:

```

sql

-- Limiting Results
-- SQL Server (T-SQL)
SELECT TOP 10 * FROM Books;

-- MySQL / PostgreSQL
SELECT * FROM Books LIMIT 10;

-- Oracle
SELECT * FROM Books WHERE ROWNUM <= 10;

-- Standard SQL (ANSI)
SELECT * FROM Books FETCH FIRST 10 ROWS ONLY;

```

1.6 Applications of SQL

1. Data Retrieval (Querying)

```

sql

-- Find all programming books under $600
SELECT Title, Author, Price
FROM Books
WHERE Genre = 'Programming' AND Price < 600;

```

2. Data Manipulation

```
sql

-- Add a new book
INSERT INTO Books (Title, Author, Price, Genre)
VALUES ('C# Programming', 'John Doe', 599.99, 'Programming');

-- Update book price
UPDATE Books SET Price = 549.99 WHERE BookID = 101;

-- Remove a book
DELETE FROM Books WHERE BookID = 101;
```

3. Database Structure Definition

```
sql

-- Create a new table
CREATE TABLE Members (
    MemberID INT PRIMARY KEY,
    Name VARCHAR(100),
    Email VARCHAR(100),
    JoinDate DATE
);

-- Modify table structure
ALTER TABLE Members ADD PhoneNumber VARCHAR(15);

-- Remove table
DROP TABLE Members;
```

4. Data Security

```
sql

-- Grant permissions
GRANT SELECT, INSERT ON Books TO LibraryStaff;

-- Revoke permissions
REVOKE DELETE ON Books FROM LibraryStaff;
```

5. Create Database Objects

sql

```
-- Create a view (virtual table)
CREATE VIEW AvailableBooks AS
SELECT Title, Author, Price
FROM Books
WHERE IsAvailable = 1;

-- Create stored procedure
CREATE PROCEDURE GetBooksByGenre
    @Genre VARCHAR(50)
AS
BEGIN
    SELECT * FROM Books WHERE Genre = @Genre;
END;

-- Create function
CREATE FUNCTION CalculateLateFee(@DaysLate INT)
RETURNS DECIMAL(10,2)
AS
BEGIN
    RETURN @DaysLate * 0.50;
END;
```

1.7 SQL Categories

SQL commands are divided into several categories:

1. DDL (Data Definition Language)

- Defines database structure
- Commands: **CREATE**, **ALTER**, **DROP**, **TRUNCATE**

sql

```
CREATE TABLE Books (BookID INT, Title VARCHAR(200));
ALTER TABLE Books ADD Author VARCHAR(100);
DROP TABLE Books;
```

2. DML (Data Manipulation Language)

- Manipulates data within tables
- Commands: **SELECT**, **INSERT**, **UPDATE**, **DELETE**

sql

```
INSERT INTO Books VALUES (101, 'C# Guide', 'John Doe');
UPDATE Books SET Price = 599.99 WHERE BookID = 101;
DELETE FROM Books WHERE BookID = 101;
```

3. DCL (Data Control Language)

- Controls access to data
- Commands: **GRANT**, **REVOKE**

sql

```
GRANT SELECT ON Books TO User1;
REVOKE INSERT ON Books FROM User1;
```

4. TCL (Transaction Control Language)

- Manages database transactions
- Commands: **COMMIT**, **ROLLBACK**, **SAVEPOINT**

sql

```
BEGIN TRANSACTION;
UPDATE Books SET Price = Price * 1.10;
COMMIT; -- or ROLLBACK;
```

1.8 Real-World Use Cases

E-Commerce Platform:

sql

```
-- Find top-selling products
SELECT ProductName, SUM(Quantity) AS TotalSold
FROM Orders
GROUP BY ProductName
ORDER BY TotalSold DESC;
```

```
-- Calculate customer lifetime value
SELECT CustomerID, SUM(OrderTotal) AS TotalSpent
FROM Orders
GROUP BY CustomerID;
```

Library Management System:

```
sql

-- Find overdue books
SELECT B.Title, M.Name, L.DueDate
FROM Loans L
JOIN Books B ON L.BookID = B.BookID
JOIN Members M ON L.MemberID = M.MemberID
WHERE L.DueDate < GETDATE() AND L.ReturnDate IS NULL;

-- Track most popular genres
SELECT Genre, COUNT(*) AS BorrowCount
FROM Loans L
JOIN Books B ON L.BookID = B.BookID
GROUP BY Genre
ORDER BY BorrowCount DESC;
```

Banking System:

```
sql

-- Check account balance
SELECT AccountNumber, Balance
FROM Accounts
WHERE CustomerID = 12345;

-- Transfer money between accounts
BEGIN TRANSACTION;
UPDATE Accounts SET Balance = Balance - 1000 WHERE AccountNumber = 'ACC001';
UPDATE Accounts SET Balance = Balance + 1000 WHERE AccountNumber = 'ACC002';
COMMIT;
```

Section 2: Understanding RDBMS

2.1 What is RDBMS?

RDBMS (Relational Database Management System) is a type of DBMS that stores data in a structured format using rows and columns, forming tables. It's based on the relational model proposed by E.F. Codd in 1970.

Key Principle: Data is organized into related tables, and relationships between tables are established through keys (Primary Keys and Foreign Keys).

2.2 The Relational Model

E.F. Codd's Relational Model (1970):

- Data is represented in relations (tables)
- Each relation consists of tuples (rows) and attributes (columns)
- Mathematical foundation based on set theory and predicate logic
- Operations performed using relational algebra

Core Concepts:

1. Relations (Tables) A relation is a two-dimensional table with:

- **Rows (Tuples):** Individual records
- **Columns (Attributes):** Data fields

2. Domains Set of allowable values for an attribute (e.g., Age domain: 0-120)

3. Keys

- **Primary Key:** Uniquely identifies each row
- **Foreign Key:** Links to primary key in another table
- **Candidate Key:** Potential primary key
- **Composite Key:** Primary key made of multiple columns

2.3 RDBMS Components Explained

2.3.1 Tables (Relations)

A **table** is a collection of related data entries consisting of rows and columns.

Structure:

Books Table:

BookID	Title	Author	Price	Genre
101	C# Programming	John Doe	599.99	Tech
102	Data Structures	Jane Smith	799.50	Tech
103	World History	Bob Johnson	699.00	History

Table Characteristics:

- **Table Name:** Unique identifier (e.g., Books, Members, Loans)
- **Schema:** Structure definition (columns, data types, constraints)
- **Data:** Actual records stored in rows

2.3.2 Columns (Fields/Attributes)

Columns are vertical entities in a table, representing a specific attribute of the data.

Example:

```
sql  
  
CREATE TABLE Members (  
    MemberID INT,          -- Column 1: Member identifier  
    Name VARCHAR(100),     -- Column 2: Member name  
    Email VARCHAR(100),    -- Column 3: Email address  
    JoinDate DATE,         -- Column 4: Registration date  
    MemberType VARCHAR(20) -- Column 5: Student/Faculty/Guest  
);
```

Column Properties:

- **Name:** Unique within the table (e.g., MemberID, Name, Email)
- **Data Type:** Defines what kind of data (INT, VARCHAR, DATE)
- **Constraints:** Rules (NOT NULL, UNIQUE, DEFAULT)
- **Size:** Maximum length for text fields (VARCHAR(100))

2.3.3 Rows (Records/Tuples)

Rows are horizontal entities representing individual data entries.

Example:

```
sql  
-- Each row is a complete record  
INSERT INTO Members VALUES (1, 'Alice Johnson', 'alice@email.com', '2024-01-15', 'Student');  
INSERT INTO Members VALUES (2, 'Bob Smith', 'bob@email.com', '2024-02-20', 'Faculty');
```

Row Characteristics:

- Each row represents one entity (one member, one book, one order)
- Order of rows doesn't matter in relational model
- Each row should be uniquely identifiable (via primary key)

2.3.4 NULL Values

NULL represents the absence of a value. It's different from:

- Zero (0) - which is a numeric value
- Empty string ("") - which is a text value
- Space (' ') - which is a character

NULL Characteristics:

```
sql  
-- NULL examples  
INSERT INTO Books (BookID, Title, Author, Publisher)  
VALUES (101, 'C# Guide', 'John Doe', NULL); -- Publisher not known yet  
  
-- NULL in comparisons  
SELECT * FROM Books WHERE Publisher IS NULL; -- ✓ Correct  
SELECT * FROM Books WHERE Publisher = NULL; -- ✗ Wrong (always false)  
  
-- NULL in calculations  
SELECT Price * Quantity AS Total FROM Orders; -- If Price is NULL, Total is NULL
```

Handling NULL:

```
sql
```

-- Check for NULL

```
SELECT * FROM Books WHERE Publisher IS NULL;
```

```
SELECT * FROM Books WHERE Publisher IS NOT NULL;
```

-- Replace NULL with default value

```
SELECT ISNULL(Publisher, 'Unknown Publisher') AS Publisher FROM Books;
```

```
SELECT COALESCE(Publisher, 'Unknown') AS Publisher FROM Books; -- Standard SQL
```

2.4 Relationships in RDBMS

Relationships define how tables are connected to each other.

Types of Relationships:

1. One-to-One (1:1) One record in Table A relates to exactly one record in Table B.

```
sql
```

-- Example: Member and MemberProfile

```
CREATE TABLE Members (
    MemberID INT PRIMARY KEY,
    Name VARCHAR(100),
    Email VARCHAR(100)
);
```

```
CREATE TABLE MemberProfiles (
    ProfileID INT PRIMARY KEY,
    MemberID INT UNIQUE, -- One-to-one constraint
    Bio TEXT,
    Photo VARBINARY(MAX),
    FOREIGN KEY (MemberID) REFERENCES Members(MemberID)
);
```

2. One-to-Many (1:N) One record in Table A relates to multiple records in Table B.

```
sql
```

-- Example: Member borrows multiple Books

CREATE TABLE Members (

MemberID **INT PRIMARY KEY**,

Name **VARCHAR(100)**

);

CREATE TABLE Loans (

LoanID **INT PRIMARY KEY**,

MemberID **INT**, -- Many loans can belong to one member

BookID **INT**,

LoanDate **DATE**,

FOREIGN KEY (MemberID) **REFERENCES** Members(MemberID)

);

3. Many-to-Many (M:N) Multiple records in Table A relate to multiple records in Table B. Requires a junction table.

sql

-- Example: Books and Authors (a book can have multiple authors, an author can write multiple books)

CREATE TABLE Books (

BookID **INT PRIMARY KEY**,

Title **VARCHAR(200)**

);

CREATE TABLE Authors (

AuthorID **INT PRIMARY KEY**,

Name **VARCHAR(100)**

);

-- Junction table

CREATE TABLE BookAuthors (

BookID **INT**,

AuthorID **INT**,

PRIMARY KEY (BookID, AuthorID),

FOREIGN KEY (BookID) **REFERENCES** Books(BookID),

FOREIGN KEY (AuthorID) **REFERENCES** Authors(AuthorID)

);

2.5 Keys in RDBMS

1. Primary Key

- Uniquely identifies each record in a table
- Cannot contain NULL values
- Only one primary key per table
- Can be single column or composite (multiple columns)

sql

```
CREATE TABLE Books (
    BookID INT PRIMARY KEY, -- Single column primary key
    Title VARCHAR(200),
    ISBN VARCHAR(20) UNIQUE
);

-- Composite primary key
CREATE TABLE BookAuthors (
    BookID INT,
    AuthorID INT,
    PRIMARY KEY (BookID, AuthorID) -- Both columns together form primary key
);
```

2. Foreign Key

- Creates relationship between tables
- References primary key in another table
- Maintains referential integrity

sql

```
CREATE TABLE Loans (
    LoanID INT PRIMARY KEY,
    MemberID INT,
    BookID INT,
    FOREIGN KEY (MemberID) REFERENCES Members(MemberID),
    FOREIGN KEY (BookID) REFERENCES Books(BookID)
);
```

3. Candidate Key

- Columns that could serve as primary key
- Unique and NOT NULL

- One candidate key is chosen as primary key

sql

```
CREATE TABLE Members (
    MemberID INT,      -- Candidate key 1
    Email VARCHAR(100), -- Candidate key 2 (unique)
    SSN VARCHAR(11),   -- Candidate key 3 (unique)
    Name VARCHAR(100),
    PRIMARY KEY (MemberID), -- Chosen as primary key
    UNIQUE (Email),
    UNIQUE (SSN)
);
```

4. Unique Key

- Ensures all values in column are different
- Can contain one NULL value
- Can have multiple unique keys in a table

sql

```
CREATE TABLE Books (
    BookID INT PRIMARY KEY,
    ISBN VARCHAR(20) UNIQUE, -- Unique key
    Title VARCHAR(200)
);
```

2.6 Constraints in RDBMS

Constraints enforce rules on data in tables to maintain data integrity.

Common Constraints:

1. PRIMARY KEY

sql

```
CREATE TABLE Books (
    BookID INT PRIMARY KEY
);
```

2. FOREIGN KEY

sql

```
CREATE TABLE Loans (
    LoanID INT PRIMARY KEY,
    BookID INT FOREIGN KEY REFERENCES Books(BookID)
);
```

3. UNIQUE

sql

```
CREATE TABLE Members (
    MemberID INT PRIMARY KEY,
    Email VARCHAR(100) UNIQUE
);
```

4. NOT NULL

sql

```
CREATE TABLE Books (
    BookID INT PRIMARY KEY,
    Title VARCHAR(200) NOT NULL, -- Title is mandatory
    Author VARCHAR(100) NOT NULL
);
```

5. CHECK

sql

```
CREATE TABLE Books (
    BookID INT PRIMARY KEY,
    Price DECIMAL(10,2) CHECK (Price >= 0), -- Price must be non-negative
    PublicationYear INT CHECK (PublicationYear BETWEEN 1900 AND 2100)
);
```

6. DEFAULT

sql

```
CREATE TABLE Members (
    MemberID INT PRIMARY KEY,
    JoinDate DATE DEFAULT GETDATE(), -- Default to current date
    Status VARCHAR(20) DEFAULT 'Active'
);
```

2.7 Advantages of RDBMS

1. Data Integrity

- Constraints ensure data accuracy
- Relationships maintain consistency
- No duplicate or orphaned records

2. Data Security

- User authentication and authorization
- Role-based access control
- Encryption support

3. Data Independence

- Physical data independence: Storage changes don't affect applications
- Logical data independence: Schema changes minimally impact applications

4. Reduced Data Redundancy

- Normalization eliminates duplication
- Data stored once, referenced many times

5. Easy Data Retrieval

- SQL provides powerful querying
- Complex queries with simple syntax
- Fast searches with indexes

6. ACID Properties

- **Atomicity:** Transactions are all-or-nothing

- **Consistency:** Database remains in valid state
- **Isolation:** Concurrent transactions don't interfere
- **Durability:** Committed data is permanent

7. Scalability

- Handles increasing data volumes
 - Supports multiple users
 - Performance optimization options
-

Section 3: Database Normalization

3.1 What is Normalization?

Database Normalization is the process of organizing data in a database to reduce redundancy and improve data integrity. It involves dividing large tables into smaller, related tables and defining relationships between them.

Goals of Normalization:

- Eliminate data redundancy (duplicate data)
- Ensure data dependencies make sense
- Reduce data anomalies
- Improve database structure
- Make database maintenance easier

Real-World Analogy: Think of organizing a messy closet. Instead of throwing everything into one big pile, you:

- Separate clothes by type (shirts, pants, shoes)
- Group similar items together
- Create logical storage locations
- Make it easy to find what you need

3.2 Data Anomalies (Problems Without Normalization)

1. Insert Anomaly

Cannot add data without other data being present.

sql

-- Poor Design: StudentCourses table

StudentID	StudentName	CourseID	CourseName
101	Alice	CS101	Programming

-- Problem: Can't add a new course unless a student enrolls in it

-- Can't add student unless they're enrolled in a course

2. Update Anomaly

Same information stored in multiple places must be updated in all locations.

sql

-- Poor Design

StudentID	StudentName	CourseID	CourseName
101	Alice	CS101	Programming
101	Alice	CS102	Databases
102	Bob	CS101	Programming

-- Problem: If course name changes from "Programming" to "Intro to Programming",

-- must update multiple rows (error-prone)

3. Delete Anomaly

Deleting one piece of information unintentionally deletes other information.

sql

-- Problem: If Alice drops CS101, we lose the information that CS101 exists

`DELETE FROM StudentCourses WHERE StudentID = 101 AND CourseID = 'CS101';`

3.3 Normal Forms

Normal forms are progressive rules. Each normal form builds upon the previous one.

Progression:

Unnormalized → 1NF → 2NF → 3NF → BCNF → 4NF → 5NF

Most databases aim for **Third Normal Form (3NF)** as it provides a good balance between normalization and performance.

3.4 First Normal Form (1NF)

Definition: A table is in 1NF if:

1. Each column contains atomic (indivisible) values
2. Each column contains values of a single type
3. Each column has a unique name
4. Order of rows doesn't matter
5. No repeating groups or arrays

Rules for 1NF:

- Atomic values only (no lists, no multiple values in one cell)
- Each cell contains single value
- Primary key defined
- No duplicate rows

Example: Violates 1NF

```
sql
```

-- **X** BAD: Not in INF

```
CREATE TABLE Members (
    MemberID INT PRIMARY KEY,
    Name VARCHAR(100),
    PhoneNumbers VARCHAR(200), -- "123-4567, 987-6543, 555-1234" (multiple values!)
    BorrowedBooks VARCHAR(500) -- "C# Guide, Data Structures, Algorithms" (list!)
);
```

-- Sample data

MemberID	Name	PhoneNumbers	BorrowedBooks
1	Alice	123-4567, 987-6543	C# Guide, Data Structures
2	Bob	555-1234	Algorithms, Web Development, SQL Basics

-- Problems:

- 1. Can't easily search for a specific phone number
- 2. Can't easily find who borrowed "C# Guide"
- 3. Can't sort by phone number
- 4. Can't establish relationships with Books table

Example: Follows 1NF

```
sql
```

-- GOOD: In INF

```
CREATE TABLE Members (
    MemberID INT PRIMARY KEY,
    Name VARCHAR(100)
);
```

-- Separate table for phone numbers (one phone per row)

```
CREATE TABLE MemberPhones (
    PhoneID INT PRIMARY KEY,
    MemberID INT,
    PhoneNumber VARCHAR(15), -- Atomic value
    PhoneType VARCHAR(20), -- Home, Mobile, Work
    FOREIGN KEY (MemberID) REFERENCES Members(MemberID)
);
```

-- Separate table for borrowed books

```
CREATE TABLE BorrowedBooks (
    LoanID INT PRIMARY KEY,
    MemberID INT,
    BookID INT,
    BorrowDate DATE,
    FOREIGN KEY (MemberID) REFERENCES Members(MemberID),
    FOREIGN KEY (BookID) REFERENCES Books(BookID)
);
```

-- Sample data

Members:

MemberID	Name
1	Alice
2	Bob

MemberPhones:

PhoneID	MemberID	PhoneNumber	PhoneType
1	1	123-4567	Home
2	1	987-6543	Mobile
3	2	555-1234	Mobile

-- Now we can easily query individual phones and books!

Converting to 1NF Steps:

1. Identify repeating groups or multi-valued attributes
 2. Create new tables for repeating data
 3. Ensure each cell contains single value
 4. Define primary keys
 5. Establish relationships with foreign keys
-

3.5 Second Normal Form (2NF)

Definition: A table is in 2NF if:

1. It is in 1NF
2. All non-key attributes are fully dependent on the entire primary key (no partial dependencies)

Partial Dependency: When a non-key column depends on only part of a composite primary key.

Key Point: 2NF only applies to tables with composite primary keys. If your table has a single-column primary key, it's automatically in 2NF if it's in 1NF.

Example: Violates 2NF

```
sql
```

-- **X** BAD: Not in 2NF (but is in 1NF)

CREATE TABLE StudentCourses (

```
StudentID INT,  
CourseID VARCHAR(10),  
StudentName VARCHAR(100), -- Depends only on StudentID (partial dependency!)  
StudentEmail VARCHAR(100), -- Depends only on StudentID (partial dependency!)  
CourseName VARCHAR(100), -- Depends only on CourseID (partial dependency!)  
Instructor VARCHAR(100), -- Depends only on CourseID (partial dependency!)  
Grade CHAR(2), -- Depends on BOTH StudentID and CourseID (full dependency) ✓)  
PRIMARY KEY (StudentID, CourseID) -- Composite primary key  
);
```

-- Sample data

StudentID	CourseID	StudentName	StudentEmail	CourseName	Instructor	Grade
101	CS101	Alice	alice@email.com	Programming	Dr. Smith	A
101	CS102	Alice	alice@email.com	Databases	Dr. Jones	B+
102	CS101	Bob	bob@email.com	Programming	Dr. Smith	A-

-- Problems:

- 1. Student information (Name, Email) repeated for each course
- 2. Course information (CourseName, Instructor) repeated for each student
- 3. Update anomaly: If Alice changes email, must update multiple rows
- 4. Delete anomaly: If student drops all courses, we lose student information

Example: Follows 2NF

sql

-- GOOD: In 2NF

-- Separate table for Students

CREATE TABLE Students (

 StudentID **INT PRIMARY KEY**,
 StudentName **VARCHAR(100)**,
 StudentEmail **VARCHAR(100)**
);

-- Separate table for Courses

CREATE TABLE Courses (

 CourseID **VARCHAR(10) PRIMARY KEY**,
 CourseName **VARCHAR(100)**,
 Instructor **VARCHAR(100)**
);

-- Junction table with only full dependencies

CREATE TABLE Enrollments (

 StudentID **INT**,
 CourseID **VARCHAR(10)**,
 Grade **CHAR(2)**,
 EnrollmentDate **DATE**,
 PRIMARY KEY (StudentID, CourseID),
 FOREIGN KEY (StudentID) **REFERENCES** Students(StudentID),
 FOREIGN KEY (CourseID) **REFERENCES** Courses(CourseID)
);

-- Sample data

Students:

StudentID	StudentName	StudentEmail
101	Alice	alice@email.com
102	Bob	bob@email.com

Courses:

CourseID	CourseName	Instructor
CS101	Programming	Dr. Smith
CS102	Databases	Dr. Jones

Enrollments:

StudentID	CourseID	Grade	EnrollmentDate
101	CS101	A	2024-01-15
101	CS102	B+	2024-01-15
102	CS101	A-	2024-01-20

-- Benefits:

- 1. No redundant student or course data
- 2. Update student email in one place only
- 3. Can add students without enrolling them in courses
- 4. Can add courses without any students enrolled

Converting to 2NF Steps:

1. Identify the primary key (especially composite keys)
2. Find columns that depend on only part of the primary key
3. Create separate tables for those partial dependencies
4. Keep only fully dependent columns in the original table

3.6 Third Normal Form (3NF)

Definition: A table is in 3NF if:

1. It is in 2NF
2. No transitive dependencies exist (non-key columns don't depend on other non-key columns)

Transitive Dependency: When a non-key column depends on another non-key column, which in turn depends on the primary key.

Pattern: PrimaryKey → ColumnA → ColumnB (ColumnB transitively depends on PrimaryKey through ColumnA)

Example: Violates 3NF

sql

-- **X** BAD: Not in 3NF (but is in 2NF)

CREATE TABLE Books (

```
BookID INT PRIMARY KEY,  
Title VARCHAR(200),  
AuthorID INT,  
AuthorName VARCHAR(100), -- Depends on AuthorID, not directly on BookID (transitive!)  
AuthorEmail VARCHAR(100), -- Depends on AuthorID, not directly on BookID (transitive!)  
AuthorCountry VARCHAR(50), -- Depends on AuthorID, not directly on BookID (transitive!)  
PublisherID INT,  
PublisherName VARCHAR(100), -- Depends on PublisherID, not directly on BookID (transitive!)  
PublisherCity VARCHAR(100), -- Depends on PublisherID, not directly on BookID (transitive!)  
Price DECIMAL(10,2),  
Genre VARCHAR(50)  
);
```

-- Sample data

BookID	Title	AuthorID	AuthorName	AuthorEmail	AuthorCountry	PublisherID	PublisherName	Publis
101	C# Programming	201	John Doe	john@email.com	USA	301	Tech Books Inc	New York
102	Data Structures	201	John Doe	john@email.com	USA	302	Code Publishers	Boston
103	World History	202	Jane Smith	jane@email.com	UK	301	Tech Books Inc	New York

-- Problems:

- 1. Author info (Name, Email, Country) repeated for each book by same author
- 2. Publisher info (Name, City) repeated for each book by same publisher
- 3. If author changes email, must update all their books
- 4. If all books by a publisher are deleted, we lose publisher information

Example: Follows 3NF

sql

-- *GOOD: In 3NF*

-- *Separate table for Books (no transitive dependencies)*

CREATE TABLE Books (

```
BookID INT PRIMARY KEY,  
Title VARCHAR(200),  
AuthorID INT,  
PublisherID INT,  
Price DECIMAL(10,2),  
Genre VARCHAR(50),  
PublicationYear INT,  
FOREIGN KEY (AuthorID) REFERENCES Authors(AuthorID),  
FOREIGN KEY (PublisherID) REFERENCES Publishers(PublisherID)  
);
```

-- *Separate table for Authors*

CREATE TABLE Authors (

```
AuthorID INT PRIMARY KEY,  
AuthorName VARCHAR(100),  
AuthorEmail VARCHAR(100),  
AuthorCountry VARCHAR(50),  
Biography TEXT
```

);

-- *Separate table for Publishers*

CREATE TABLE Publishers (

```
PublisherID INT PRIMARY KEY,  
PublisherName VARCHAR(100),  
PublisherCity VARCHAR(100),  
PublisherCountry VARCHAR(50),  
Website VARCHAR(200)
```

);

-- *Sample data*

Books:

BookID	Title	AuthorID	PublisherID	Price	Genre	PublicationYear
101	C# Programming	201	301	599.99	Programming	2024
102	Data Structures	201	302	799.50	Programming	2023
103	World History	202	301	699.00	History	2022

Authors:

AuthorID	AuthorName	AuthorEmail	AuthorCountry
201	John Doe	john@email.com	USA
202	Jane Smith	jane@email.com	UK

Publishers:

PublisherID	PublisherName	PublisherCity
301	Tech Books Inc	New York
302	Code Publishers	Boston

-- Benefits:

- 1. Author information stored once, referenced by BookID
- 2. Publisher information stored once, referenced by BookID
- 3. Update author email in one place
- 4. Can add authors without books
- 5. Can add publishers without books
- 6. No redundant data

Converting to 3NF Steps:

1. Identify non-key columns that depend on other non-key columns
2. Create separate tables for those columns
3. Use foreign keys to establish relationships
4. Remove transitive dependencies from original table

3.7 Normalization Benefits & Trade-offs

Benefits of Normalization:

1. Eliminates Data Redundancy

- Data stored once, referenced many times
- Reduces storage requirements
- Consistent data across database

2. Prevents Data Anomalies

- No insert anomalies
- No update anomalies
- No delete anomalies

3. Improves Data Integrity

- Constraints ensure accuracy
- Relationships maintain consistency
- Easier to maintain valid data

4. Easier Maintenance

- Changes in one place
- Less code to update
- Fewer errors

5. Flexible Schema

- Easy to add new entities
- Easy to modify relationships
- Scalable structure

Trade-offs (Denormalization Considerations):

1. More Complex Queries

- Need JOINS to retrieve data
- More tables to manage
- Longer query execution time

2. Performance Impact

- JOINS can be expensive
- More disk I/O for related data
- May need indexes for performance

When to Denormalize:

- Read-heavy applications (reporting, analytics)
- Performance is critical
- Data rarely changes
- Complex queries are common

Example: Denormalization for Performance

```
sql

-- Normalized (3NF)
SELECT B.Title, A.AuthorName, P.PublisherName
FROM Books B
JOIN Authors A ON B.AuthorID = A.AuthorID
JOIN Publishers P ON B.PublisherID = P.PublisherID;

-- Denormalized (for read performance)
CREATE TABLE BooksWithDetails (
    BookID INT PRIMARY KEY,
    Title VARCHAR(200),
    AuthorName VARCHAR(100), -- Denormalized
    PublisherName VARCHAR(100), -- Denormalized
    Price DECIMAL(10,2)
);

-- Faster query (no JOINs)
SELECT Title, AuthorName, PublisherName FROM BooksWithDetails;
```

Section 4: SQL Data Types

4.1 Why Data Types Matter

Data Types define the kind of data that can be stored in a column. Choosing the right data type is crucial for:

1. Data Integrity

- Ensures only valid data is stored
- Prevents errors (e.g., storing text in numeric field)
- Enforces business rules

2. Storage Optimization

- Uses appropriate space for each value
- Reduces database size
- Improves performance

3. Query Performance

- Faster searches and comparisons
- Better index performance
- Optimized joins

4. Data Validation

- Automatic type checking
- Prevents invalid data entry
- Reduces application-level validation

Example:

```
sql
```

--  Poor Choice: Using VARCHAR for numbers

CREATE TABLE Products (

```
ProductID VARCHAR(50), -- Should be INT  
Price VARCHAR(20), -- Should be DECIMAL  
Quantity VARCHAR(10) -- Should be INT  
);
```

-- Problems:

- Can store "ABC" in ProductID
- Can store "expensive" in Price
- Can't perform mathematical calculations easily
- Takes more space than needed
- Slower comparisons and sorting

--  Good Choice: Appropriate data types

CREATE TABLE Products (

```
ProductID INT PRIMARY KEY,  
Price DECIMAL(10,2),  
Quantity INT  
);
```

-- Benefits:

- Only numbers allowed
- Precise decimal calculations
- Efficient storage
- Fast operations

4.2 String Data Types

String data types store text, characters, and symbols.

4.2.1 CHAR(n) - Fixed-Length Character String

Characteristics:

- Fixed length specified by \boxed{n}
- Range: 1 to 8,000 characters
- Pads with spaces if data is shorter than \boxed{n}
- Always uses \boxed{n} bytes of storage
- Faster for fixed-length data

When to Use:

- Data is always the same length (e.g., country codes, state codes)
- Performance is critical for fixed-length strings
- Storage predictability is important

Examples:

```
sql
```

```
CREATE TABLE Countries (
    CountryCode CHAR(2),      -- 'US', 'UK', 'IN' - always 2 characters
    CurrencyCode CHAR(3),     -- 'USD', 'GBP', 'INR' - always 3 characters
    PostalCode CHAR(6)        -- Fixed-length postal codes
);
```

```
INSERT INTO Countries VALUES ('US', 'USD', '123456');
INSERT INTO Countries VALUES ('UK', 'GBP', '789012');
```

```
-- Storage: Each CHAR(2) uses 2 bytes, even if you store 'A'
-- 'A' is stored as 'A ' (with trailing space)
```

Comparison: CHAR vs VARCHAR

```
sql

-- CHAR(10) storage
'Hello'  stored as 'Hello   ' (10 bytes - padded with spaces)
'Database' stored as 'Database ' (10 bytes - padded with spaces)

-- VARCHAR(10) storage
'Hello'  stored as 'Hello' (5 bytes + 2 bytes overhead)
'Database' stored as 'Database' (8 bytes + 2 bytes overhead)
```

4.2.2 VARCHAR(n) - Variable-Length Character String

Characteristics:

- Variable length up to $\lceil n \rceil$ characters
- Range: 1 to 8,000 characters
- Uses actual length + 2 bytes overhead

- No padding
- Most commonly used string type

When to Use:

- Data length varies (names, emails, addresses)
- Want to save storage space
- Most general-purpose text storage

Examples:

sql

```
CREATE TABLE Members (
    MemberID INT PRIMARY KEY,
    FirstName VARCHAR(50),      -- Variable length
    LastName VARCHAR(50),
    Email VARCHAR(100),         -- Emails vary in length
    Address VARCHAR(200),        -- Addresses vary significantly
    City VARCHAR(50)
);

INSERT INTO Members VALUES
(1, 'John', 'Doe', 'john.doe@email.com', '123 Main St, Apt 4B', 'New York'),
(2, 'Alice', 'Johnson-Smith', 'alice.j.smith@company.com', '4567 Oak Avenue', 'Boston');

-- Storage:
-- 'John' uses 4 bytes + 2 bytes overhead = 6 bytes total
-- 'Alice' uses 5 bytes + 2 bytes overhead = 7 bytes total
```

4.2.3 VARCHAR(MAX) - Large Variable-Length Text

Characteristics:

- Variable length up to 2GB ($2^{31}-1$ bytes)
- Replaces deprecated TEXT data type
- Stored outside main data page if exceeds 8KB
- Use for large text that varies significantly

When to Use:

- Blog posts, articles, descriptions
- Long comments or reviews
- Document content
- Data with unpredictable large size

Examples:

sql

```
CREATE TABLE BlogPosts (
    PostID INT PRIMARY KEY,
    Title VARCHAR(200),
    Content VARCHAR(MAX),      -- Can store entire blog post
    Summary VARCHAR(500)
);
```

```
CREATE TABLE Books (
    BookID INT PRIMARY KEY,
    Title VARCHAR(200),
    Description VARCHAR(MAX),  -- Can store full book description
    FullText VARCHAR(MAX)      -- Can store entire book content
);
```

4.2.4 NCHAR(n) & NVARCHAR(n) - Unicode Strings

Characteristics:

- **NCHAR(n):** Fixed-length Unicode (2 bytes per character)
- **NVARCHAR(n):** Variable-length Unicode (up to 4,000 characters)
- **NVARCHAR(MAX):** Variable-length Unicode (up to 2GB)
- Supports international characters (Chinese, Arabic, Emoji, etc.)
- Uses twice the storage of non-Unicode types

When to Use:

- International applications
- Multiple language support
- Special characters, symbols, emoji
- When you need to store characters from any language

Examples:

sql

```
CREATE TABLE InternationalMembers (
    MemberID INT PRIMARY KEY,
    Name NVARCHAR(100),          -- Can store names in any language
    Bio NVARCHAR(MAX),           -- Can store text in any language
    Country NVARCHAR(50)
);

-- Can store international characters
INSERT INTO InternationalMembers VALUES
(1, N'李明', N'来自中国的软件工程师', N'中国'), -- Chinese
(2, N'محمد أحمد', N'مطور برمجيات من مصر'), -- Arabic
(3, N'Анна Иванова', N'Разработчик из России', N'Россия'); -- Russian

-- Note: Use N prefix for Unicode strings: N'text'
```

String Data Types Summary:

Data Type	Max Length	Storage	Use Case	Example
CHAR(n)	8,000	Fixed n bytes	Fixed-length codes	Country code: 'US'
VARCHAR(n)	8,000	Variable + 2 bytes	Variable text	Names, emails
VARCHAR(MAX)	2GB	Variable	Large text	Articles, descriptions
NCHAR(n)	4,000	Fixed n×2 bytes	Fixed Unicode	Fixed international codes
NVARCHAR(n)	4,000	Variable + 2 bytes	Variable Unicode	International names
NVARCHAR(MAX)	2GB	Variable	Large Unicode text	Multilingual content

4.3 Numeric Data Types

Numeric data types store numbers - integers, decimals, and floating-point values.

4.3.1 Integer Types

TINYINT

- Range: 0 to 255
- Storage: 1 byte
- Use for: Very small numbers (age, percentage, status codes)

```
sql
```

```
CREATE TABLE Users (
    UserID INT PRIMARY KEY,
    Age TINYINT,          -- 0-255 is enough for age
    StatusCode TINYINT     -- Small status codes
);
```

SMALLINT

- Range: -32,768 to 32,767
- Storage: 2 bytes
- Use for: Small numbers (quantity, year offset)

```
sql
```

```
CREATE TABLE Products (
    ProductID INT PRIMARY KEY,
    StockQuantity SMALLINT,   -- -32K to 32K is usually enough
    ReorderLevel SMALLINT
);
```

INT (Most Common)

- Range: -2,147,483,648 to 2,147,483,647 (~2.1 billion)
- Storage: 4 bytes
- Use for: General-purpose integer storage (IDs, counts, quantities)

```
sql
```

```

CREATE TABLE Orders (
    OrderID INT PRIMARY KEY,
    CustomerID INT,
    TotalItems INT,
    OrderYear INT      -- Year: 2024
);

```

BIGINT

- Range: -9,223,372,036,854,775,808 to 9,223,372,036,854,775,807
- Storage: 8 bytes
- Use for: Very large numbers (global IDs, large counts, timestamps)

sql

```

CREATE TABLE Analytics (
    EventID BIGINT PRIMARY KEY,    -- Billions of events
    Timestamp BIGINT,            -- Unix timestamp in milliseconds
    ViewCount BIGINT             -- Can exceed 2 billion
);

```

Integer Types Comparison:

Type	Range	Storage	Use Case
TINYINT	0 to 255	1 byte	Age, small counts, flags
SMALLINT	-32,768 to 32,767	2 bytes	Small quantities, years
INT	-2.1B to 2.1B	4 bytes	General IDs, counts (most common)
BIGINT	-9.2 quintillion to 9.2 quintillion	8 bytes	Very large numbers, timestamps

4.3.2 DECIMAL & NUMERIC Types (Fixed-Point)

DECIMAL(p, s) / NUMERIC(p, s)

- **p (precision):** Total number of digits (1-38)
- **s (scale):** Number of digits after decimal point (0 to p)
- Storage: 5-17 bytes (depends on precision)

- **Exact numeric values** - no rounding errors
- **Always use for money!**

When to Use:

- Money, currency, prices
- Financial calculations
- Any calculation requiring exact precision
- Percentages with decimal places

Examples:

sql

```
CREATE TABLE Products (
    ProductID INT PRIMARY KEY,
    Price DECIMAL(10, 2),      -- 99,999,999.99 (10 digits, 2 after decimal)
    Discount DECIMAL(5, 2),    -- 999.99 (percentage like 15.50%)
    Weight DECIMAL(8, 3),    -- 99,999.999 (3 decimal places)
    TaxRate DECIMAL(5, 4)     -- 9.9999 (like 0.0875 for 8.75%)
);

-- Financial calculations with DECIMAL
INSERT INTO Products VALUES (101, 599.99, 10.00, 2.450, 0.0875);
```

```
-- Calculate total with tax (no rounding errors!)
SELECT
    ProductID,
    Price,
    Price * (1 - Discount/100) AS DiscountedPrice,
    Price * (1 - Discount/100) * (1 + TaxRate) AS FinalPrice
FROM Products;
```

Common DECIMAL Sizes:

sql

```
DECIMAL(10, 2) -- Money: $99,999,999.99
DECIMAL(18, 2) -- Large money: $9,999,999,999,999,999.99
DECIMAL(5, 2) -- Percentages: 999.99%
DECIMAL(8, 3) -- Precise measurements: 99,999.999
```

4.3.3 Money Types

MONEY

- Range: -922,337,203,685,477.5808 to 922,337,203,685,477.5807
- Storage: 8 bytes
- Precision: 4 decimal places
- Currency-optimized

SMALLMONEY

- Range: -214,748.3648 to 214,748.3647
- Storage: 4 bytes
- Precision: 4 decimal places

Examples:

```
sql
```

```
CREATE TABLE Transactions (
    TransactionID INT PRIMARY KEY,
    Amount MONEY,           -- For most financial transactions
    Fee SMALLMONEY,         -- For small fees
    Balance MONEY
);

-- Note: Modern practice prefers DECIMAL(18,2) over MONEY for better portability
```

4.3.4 Floating-Point Types (Approximate)

FLOAT(n)

- Approximate numeric value
- **n:** Number of bits (1-53)
- FLOAT(24) = 4 bytes (REAL)
- FLOAT(53) = 8 bytes (default FLOAT)
- **Use for:** Scientific calculations, statistics

REAL

- Equivalent to FLOAT(24)
- Storage: 4 bytes
- Precision: ~7 digits

When to Use:

- Scientific calculations
- Statistical analysis
- When approximate values are acceptable
- **Never use for money!**

Examples:

```
sql

CREATE TABLE ScienceData (
    ExperimentID INT PRIMARY KEY,
    Temperature FLOAT,          -- Scientific measurements
    Pressure FLOAT,
    Concentration REAL,
    Latitude FLOAT,            -- Geographic coordinates
    Longitude FLOAT
);

-- Floating-point example
INSERT INTO ScienceData VALUES
(1, 98.6, 101.325, 0.00123, 40.7128, -74.0060);
```

⚠ WARNING: Float Precision Issues

```
sql

-- Floating-point arithmetic can have rounding errors
DECLARE @f1 FLOAT = 0.1;
DECLARE @f2 FLOAT = 0.2;
SELECT @f1 + @f2; -- Might not exactly equal 0.3!

-- Always use DECIMAL for exact calculations
DECLARE @d1 DECIMAL(10,2) = 0.10;
DECLARE @d2 DECIMAL(10,2) = 0.20;
SELECT @d1 + @d2; -- Exactly 0.30
```

4.3.5 BIT (Boolean)

BIT

- Stores 0 or 1 (or NULL)
- Storage: 1 bit (but stored as byte)
- Use for: Boolean flags (Yes/No, True/False, Active/Inactive)

Examples:

```
sql
```

```
CREATE TABLE Books (
    BookID INT PRIMARY KEY,
    Title VARCHAR(200),
    IsAvailable BIT,          -- 1 = Available, 0 = Not Available
    IsEBook BIT,              -- 1 = EBook, 0 = Physical
    IsBestseller BIT,         -- 1 = Yes, 0 = No
    IsDeleted BIT DEFAULT 0  -- Soft delete flag
);

INSERT INTO Books VALUES
(101, 'C# Programming', 1, 0, 1, 0); -- Available, Physical, Bestseller, Not Deleted

-- Query available books
SELECT * FROM Books WHERE IsAvailable = 1;

-- Common practice: Use meaningful column names
IsActive, isEnabled, HasPermission, CanEdit, ShouldNotify
```

Numeric Data Types Summary:

Type	Range	Storage	Precision	Use Case
TINYINT	0-255	1 byte	Exact	Age, small counts
SMALLINT	-32K to 32K	2 bytes	Exact	Small numbers
INT	-2.1B to 2.1B	4 bytes	Exact	General purpose
BIGINT	Very large	8 bytes	Exact	Large numbers
DECIMAL(p,s)	Defined	5-17 bytes	Exact	Money (always!) , precise calculations
MONEY	Large	8 bytes	4 decimals	Currency
FLOAT	Approximate	8 bytes	~15 digits	Scientific calculations
REAL	Approximate	4 bytes	~7 digits	Scientific calculations
BIT	0/1	1 bit	-	Boolean flags

4.4 Date and Time Data Types

Date and time types store temporal data - dates, times, and timestamps.

4.4.1 DATE

Characteristics:

- Stores only date (no time component)
- Range: January 1, 0001 to December 31, 9999
- Storage: 3 bytes
- Format: 'YYYY-MM-DD'

When to Use:

- Birth dates
- Start/end dates
- Deadlines
- Any date without time component

Examples:

sql

```
CREATE TABLE Members (
    MemberID INT PRIMARY KEY,
    Name VARCHAR(100),
    BirthDate DATE,          -- Only date needed
    JoinDate DATE,
    MembershipExpiry DATE
);

INSERT INTO Members VALUES
(1, 'Alice', '1995-03-15', '2024-01-10', '2025-01-10');

-- Query members born in 1995
SELECT * FROM Members WHERE YEAR(BirthDate) = 1995;

-- Calculate age
SELECT
    Name,
    BirthDate,
    DATEDIFF(YEAR, BirthDate, GETDATE()) AS Age
FROM Members;
```

4.4.2 TIME

Characteristics:

- Stores only time (no date component)
- Range: 00:00:00.0000000 to 23:59:59.9999999
- Storage: 3-5 bytes (depends on precision)
- Format: 'HH:MM:SS.nnnnnnn'

When to Use:

- Opening/closing times
- Event duration
- Appointment times (when date is separate)
- Time of day

Examples:

sql

```
CREATE TABLE LibraryHours (
    DayOfWeek VARCHAR(10),
    OpenTime TIME,          -- Only time needed
    CloseTime TIME
);

INSERT INTO LibraryHours VALUES
('Monday', '09:00:00', '21:00:00'),
('Saturday', '10:00:00', '18:00:00');

-- Check if library is currently open
DECLARE @CurrentTime TIME = CAST(GETDATE() AS TIME);
SELECT * FROM LibraryHours
WHERE @CurrentTime BETWEEN OpenTime AND CloseTime;
```

4.4.3 DATETIME

Characteristics:

- Stores both date and time
- Range: January 1, 1753 to December 31, 9999
- Storage: 8 bytes
- Precision: 0.003 seconds (rounded to .000, .003, or .007)
- Format: 'YYYY-MM-DD HH:MM:SS.mmm'

When to Use:

- Transaction timestamps
- Log entries
- Created/modified dates with time
- Legacy systems (use DATETIME2 for new systems)

Examples:

sql

```

CREATE TABLE Transactions (
    TransactionID INT PRIMARY KEY,
    Amount DECIMAL(10,2),
    TransactionDate DATETIME DEFAULT GETDATE() -- Current date and time
);

INSERT INTO Transactions (TransactionID, Amount)
VALUES (1, 599.99); -- TransactionDate automatically set to current datetime

-- Query transactions from last 7 days
SELECT * FROM Transactions
WHERE TransactionDate >= DATEADD(DAY, -7, GETDATE());

```

4.4.4 DATETIME2

Characteristics:

- Stores both date and time (improved DATETIME)
- Range: January 1, 0001 to December 31, 9999
- Storage: 6-8 bytes
- Precision: 100 nanoseconds (7 fractional seconds by default)
- Format: 'YYYY-MM-DD HH:MM:SS.nnnnnnnn'
- **Recommended over DATETIME for new applications**

When to Use:

- Prefer over DATETIME for all new applications
- When you need greater precision
- When you need wider date range
- Modern timestamp requirements

Examples:

sql

```

CREATE TABLE AuditLog (
    LogID INT PRIMARY KEY,
    Action VARCHAR(100),
    LogTimestamp DATETIME2 DEFAULT SYSDATETIME(), -- High precision
    UserID INT
);

-- Specify precision
CREATE TABLE Events (
    EventID INT PRIMARY KEY,
    EventTime DATETIME2(3) -- 3 fractional seconds: .123
);

```

4.4.5 SMALLDATETIME

Characteristics:

- Stores both date and time (smaller range)
- Range: January 1, 1900 to June 6, 2079
- Storage: 4 bytes
- Precision: 1 minute
- Format: 'YYYY-MM-DD HH:MM:00'

When to Use:

- When you need to save space
- Minute-level precision is sufficient
- Date range is within limits

Examples:

sql

```
CREATE TABLE Appointments (
    AppointmentID INT PRIMARY KEY,
    PatientName VARCHAR(100),
    AppointmentTime SMALLDATETIME -- Minute precision is enough
);
```

```
INSERT INTO Appointments VALUES
(1, 'John Doe', '2024-12-20 14:30'); -- Seconds are ignored
```

4.4.6 DATETIMEOFFSET

Characteristics:

- Stores date, time, AND time zone offset
- Range: January 1, 0001 to December 31, 9999
- Storage: 10 bytes
- Precision: 100 nanoseconds
- Format: 'YYYY-MM-DD HH:MM:SS.nnnnnnnn +/-HH:MM'

When to Use:

- International applications
- Need to store time zone information
- Scheduling across time zones
- Audit trails with time zone

Examples:

sql

```

CREATE TABLE GlobalEvents (
    EventID INT PRIMARY KEY,
    EventName VARCHAR(100),
    EventTime DATETIMEOFFSET -- Stores time zone offset
);

-- Insert with time zone
INSERT INTO GlobalEvents VALUES
(1, 'New York Meeting', '2024-12-20 09:00:00 -05:00'), -- EST
(2, 'London Meeting', '2024-12-20 14:00:00 +00:00'); -- UTC

-- Convert to different time zones
SELECT
    EventName,
    EventTime AS OriginalTime,
    SWITCHOFFSET(EventTime, '+00:00') AS UTC_Time,
    SWITCHOFFSET(EventTime, '+05:30') AS India_Time
FROM GlobalEvents;

```

4.4.7 TIMESTAMP (ROWVERSION)

Characteristics:

- Automatically generated binary number
- Unique within database
- Changes every time row is modified
- Storage: 8 bytes
- **Not a date/time type!** (despite the name)

When to Use:

- Optimistic concurrency control
- Detect row changes
- Version tracking

Examples:

sql

```

CREATE TABLE Products (
    ProductID INT PRIMARY KEY,
    ProductName VARCHAR(100),
    Price DECIMAL(10,2),
    RowVersion TIMESTAMP      -- Automatically managed
);

-- Optimistic concurrency check
DECLARE @OriginalVersion BINARY(8);

-- Read product and its version
SELECT @OriginalVersion = RowVersion FROM Products WHERE ProductID = 101;

-- Update only if version hasn't changed
UPDATE Products
SET Price = 599.99
WHERE ProductID = 101 AND RowVersion = @OriginalVersion;

IF @@ROWCOUNT = 0
BEGIN
    PRINT 'Product was modified by another user!';
END

```

Date/Time Data Types Summary:

Type	Date	Time	Range	Storage	Precision	Use Case
DATE	✓	✗	0001-9999	3 bytes	1 day	Birth dates, deadlines
TIME	✗	✓	Full day	3-5 bytes	100 ns	Opening hours, durations
DATETIME	✓	✓	1753-9999	8 bytes	3.33 ms	Legacy timestamps
DATETIME2	✓	✓	0001-9999	6-8 bytes	100 ns	Modern timestamps (preferred)
SMALLDATETIME	✓	✓	1900-2079	4 bytes	1 minute	Space-saving timestamps
DATETIMEOFFSET	✓	✓	0001-9999	10 bytes	100 ns	Time zone-aware
TIMESTAMP	✗	✗	-	8 bytes	-	Row versioning

Common Date/Time Functions:

```

sql

-- Current date/time
SELECT GETDATE();          -- DATETIME (current datetime)
SELECT SYSDATETIME();       -- DATETIME2 (high precision)
SELECT CURRENT_TIMESTAMP;   -- DATETIME (ANSI standard)

-- Date parts
SELECT YEAR(GETDATE());    -- 2024
SELECT MONTH(GETDATE());    -- 12
SELECT DAY(GETDATE());      -- 17

-- Date arithmetic
SELECT DATEADD(DAY, 7, GETDATE());    -- Add 7 days
SELECT DATEADD(MONTH, -3, GETDATE());   -- Subtract 3 months
SELECT DATEDIFF(DAY, '2024-01-01', GETDATE()); -- Days between dates

-- Formatting
SELECT FORMAT(GETDATE(), 'yyyy-MM-dd');    -- '2024-12-17'
SELECT FORMAT(GETDATE(), 'MMMM dd, yyyy');    -- 'December 17, 2024'
SELECT CONVERT(VARCHAR, GETDATE(), 101);        -- '12/17/2024'

```

4.5 Other Important Data Types

4.5.1 UNIQUEIDENTIFIER (GUID)

Characteristics:

- Stores Globally Unique Identifier (GUID)
- Format: 'XXXXXXXX-XXXX-XXXX-XXXX-XXXXXXXXXXXX' (32 hex digits + 4 hyphens)
- Storage: 16 bytes
- Unique across tables, databases, servers

When to Use:

- Distributed systems
- Replication scenarios
- Need globally unique IDs
- Merging data from multiple sources

Examples:

```
sql

CREATE TABLE DistributedUsers (
    UserGUID UNIQUEIDENTIFIER PRIMARY KEY DEFAULT NEWID(),
    UserName VARCHAR(100)
);

-- Insert with auto-generated GUID
INSERT INTO DistributedUsers (UserName) VALUES ('Alice');

-- Insert with specific GUID
INSERT INTO DistributedUsers VALUES
('550e8400-e29b-41d4-a716-446655440000', 'Bob');

-- Generate new GUID
SELECT NEWID(); -- e.g., '6F9619FF-8B86-D011-B42D-00C04FC964FF'
```

4.5.2 XML

Characteristics:

- Stores XML data
- Maximum size: 2GB
- Built-in XML validation and querying
- Storage: Variable

When to Use:

- Storing structured documents
- Configuration files
- SOAP messages
- Hierarchical data

Examples:

```
sql
```

```

CREATE TABLE Products (
    ProductID INT PRIMARY KEY,
    ProductName VARCHAR(100),
    Specifications XML
);

-- Insert XML data
INSERT INTO Products VALUES (
    101,
    'Laptop',
    '<specs>
        <processor>Intel i7</processor>
        <ram>16GB</ram>
        <storage>512GB SSD</storage>
    </specs>'
);

-- Query XML data
SELECT
    ProductName,
    Specifications.value('(/specs/processor)[1]', 'VARCHAR(50)') AS Processor,
    Specifications.value('(/specs/ram)[1]', 'VARCHAR(20)') AS RAM
FROM Products;

```

4.5.3 Binary Data Types

BINARY(n)

- Fixed-length binary data
- Range: 1 to 8,000 bytes
- Storage: n bytes

VARBINARY(n)

- Variable-length binary data
- Range: 1 to 8,000 bytes
- Storage: Actual length + 2 bytes

VARBINARY(MAX)

- Variable-length binary data

- Maximum: 2GB
- Storage: Actual length

When to Use:

- Store files (images, PDFs, documents)
- Store encrypted data
- Store binary serialized objects

Examples:

```
sql

CREATE TABLE Documents (
    DocumentID INT PRIMARY KEY,
    FileName VARCHAR(200),
    FileType VARCHAR(50),
    FileContent VARBINARY(MAX), -- Store actual file
    UploadDate DATETIME2 DEFAULT SYSDATETIME()
);
```

-- Note: For large files, consider storing file path instead

```
CREATE TABLE DocumentsWithPath (
    DocumentID INT PRIMARY KEY,
    FileName VARCHAR(200),
    FilePath VARCHAR(500), -- Store path instead of content
    UploadDate DATETIME2 DEFAULT SYSDATETIME()
);
```

4.5.4 JSON (Stored as NVARCHAR)

Note: SQL Server doesn't have a native JSON data type, but provides JSON functions for NVARCHAR data.

Examples:

```
sql
```

```

CREATE TABLE Products (
    ProductID INT PRIMARY KEY,
    ProductName VARCHAR(100),
    Attributes NVARCHAR(MAX)      -- Store JSON
);

-- Insert JSON data
INSERT INTO Products VALUES (
    101,
    'Laptop',
    '{"brand": "Dell", "processor": "Intel i7", "ram": "16GB", "price": 999.99}'
);

-- Query JSON data
SELECT
    ProductName,
    JSON_VALUE(Attributes, '$.brand') AS Brand,
    JSON_VALUE(Attributes, '$.price') AS Price
FROM Products;

-- Check if JSON is valid
SELECT ISJSON(Attributes) FROM Products;

```

Section 5: Practical Examples & Best Practices

5.1 Complete Library Management System - Data Types Example

sql

-- Members table

```
CREATE TABLE Members (
    MemberID INT PRIMARY KEY IDENTITY(1,1),          -- Auto-incrementing ID
    MemberGUID UNIQUEIDENTIFIER DEFAULT NEWID(),      -- Global unique ID
    FirstName NVARCHAR(50) NOT NULL,                  -- Unicode for international names
    LastName NVARCHAR(50) NOT NULL,
    Email VARCHAR(100) UNIQUE NOT NULL,
    PhoneNumber VARCHAR(15),
    BirthDate DATE,                                 -- Date only
    JoinDate DATETIME2 DEFAULT SYSDATETIME(),        -- Timestamp with precision
    MembershipExpiry DATE,
    Address NVARCHAR(200),
    City NVARCHAR(50),
    PostalCode VARCHAR(10),
    CountryCode CHAR(2),                            -- Fixed 2 characters: 'US', 'UK'
    IsActive BIT DEFAULT 1,                         -- Boolean flag
    MemberType VARCHAR(20) CHECK (MemberType IN ('Student', 'Faculty', 'Guest')),
    ProfilePhoto VARBINARY(MAX),                    -- Binary data for photo
    Preferences NVARCHAR(MAX),                     -- JSON preferences
    CreatedDate DATETIME2 DEFAULT SYSDATETIME(),
    ModifiedDate DATETIME2,
    RowVersion ROWVERSION                         -- Concurrency control
);
```

-- Books table

```
CREATE TABLE Books (
    BookID INT PRIMARY KEY IDENTITY(1,1),
    ISBN VARCHAR(20) UNIQUE NOT NULL,
    Title NVARCHAR(200) NOT NULL,
    Subtitle NVARCHAR(200),
    Author NVARCHAR(100) NOT NULL,
    Publisher NVARCHAR(100),
    PublicationDate DATE,
    Edition TINYINT,                                -- 1-255 is enough
    Pages SMALLINT,                                 -- Max ~32K pages
    Price DECIMAL(10,2) NOT NULL,                   -- Always DECIMAL for money!
    CostPrice DECIMAL(10,2),
    Genre VARCHAR(50),
    Language CHAR(3),                             -- ISO language code: 'ENG', 'SPA'
    Description NVARCHAR(MAX),
    CoverImage VARBINARY(MAX),
    ISBN13 VARCHAR(17),
    ISBN10 VARCHAR(13),
```

```
Quantity SMALLINT DEFAULT 0,  
AvailableQuantity SMALLINT DEFAULT 0,  
IsEBook BIT DEFAULT 0,  
IsBestseller BIT DEFAULT 0,  
IsActive BIT DEFAULT 1,  
AverageRating DECIMAL(3,2),  
-- 0.00 to 9.99  
TotalReviews INT DEFAULT 0,  
AddedDate DATETIME2 DEFAULT SYSDATETIME(),  
LastModified DATETIME2  
);
```

-- Loans table

```
CREATE TABLE Loans (  
    LoanID INT PRIMARY KEY IDENTITY(1,1),  
    MemberID INT NOT NULL,  
    BookID INT NOT NULL,  
    LoanDate DATETIME2 DEFAULT SYSDATETIME(),  
    DueDate DATE NOT NULL,  
    ReturnDate DATETIME2,  
    LateFee DECIMAL(10,2) DEFAULT 0.00,  
    LateFeePaid BIT DEFAULT 0,  
    Status VARCHAR(20) DEFAULT 'Active' CHECK (Status IN ('Active', 'Returned', 'Overdue', 'Lost')),  
    RenewalCount TINYINT DEFAULT 0,  
    Notes NVARCHAR(500),  
    FOREIGN KEY (MemberID) REFERENCES Members(MemberID),  
    FOREIGN KEY (BookID) REFERENCES Books(BookID)  
);
```

-- Fines table

```
CREATE TABLE Fines (  
    FineID INT PRIMARY KEY IDENTITY(1,1),  
    LoanID INT NOT NULL,  
    MemberID INT NOT NULL,  
    FineAmount DECIMAL(10,2) NOT NULL,  
    FineDate DATE DEFAULT CAST(SYSDATETIME() AS DATE),  
    DaysOverdue SMALLINT,  
    PaidAmount DECIMAL(10,2) DEFAULT 0.00,  
    PaidDate DATETIME2,  
    IsPaid BIT DEFAULT 0,  
    PaymentMethod VARCHAR(20),  
    Notes NVARCHAR(200),  
    FOREIGN KEY (LoanID) REFERENCES Loans(LoanID),  
    FOREIGN KEY (MemberID) REFERENCES Members(MemberID)  
);
```

```
-- Library Hours table
CREATE TABLE LibraryHours (
    DayOfWeek VARCHAR(10) PRIMARY KEY,
    OpenTime TIME NOT NULL,
    CloseTime TIME NOT NULL,
    IsOpen BIT DEFAULT 1
);

-- Reservations table
CREATE TABLE Reservations (
    ReservationID INT PRIMARY KEY IDENTITY(1,1),
    MemberID INT NOT NULL,
    BookID INT NOT NULL,
    ReservationDate DATETIME2 DEFAULT SYSDATETIME(),
    ExpiryDate DATETIME2,
    Status VARCHAR(20) DEFAULT 'Pending',
    NotificationSent BIT DEFAULT 0,
    FOREIGN KEY (MemberID) REFERENCES Members(MemberID),
    FOREIGN KEY (BookID) REFERENCES Books(BookID)
);
```

5.2 Data Type Selection Best Practices

1. Choose the Smallest Appropriate Data Type

```
sql

--  Wasteful
CREATE TABLE Users (
    Age BIGINT,          -- BIGINT for age? Overkill!
    Name VARCHAR(1000)   -- 1000 characters for name? Too much!
);

--  Efficient
CREATE TABLE Users (
    Age TINYINT,         -- 0-255 is perfect for age
    Name VARCHAR(100)    -- 100 characters is reasonable
);
```

2. Always Use DECIMAL for Money

```
sql
```

-- Wrong - will have rounding errors

CREATE TABLE Products (

Price **FLOAT**

);

-- Correct - exact precision

CREATE TABLE Products (

Price **DECIMAL(10,2)** -- Always for money!

);

3. Use Appropriate String Types

sql

-- Inefficient

CREATE TABLE Countries (

CountryCode **VARCHAR(100)**, -- Too large for 2-character code

Name **VARCHAR(50)** -- May truncate longer names

);

-- Efficient

CREATE TABLE Countries (

CountryCode **CHAR(2)**, -- Fixed 2 characters

Name **NVARCHAR(100)** -- Variable, international

);

4. Use Modern Date/Time Types

sql

-- Old way

CREATE TABLE Events (

EventTime **DATETIME**

);

-- Modern way

CREATE TABLE Events (

EventTime **DATETIME2** -- Better precision and range

);

5. Use BIT for Boolean Flags

sql

-- Wasteful

CREATE TABLE Users (

```
IsActive VARCHAR(10),      -- Storing 'True'/'False'?  
HasPermission INT          -- Using INT for yes/no?  
);
```

-- Efficient

CREATE TABLE Users (

```
IsActive BIT,            -- 1 or 0  
HasPermission BIT  
);
```

5.3 Common Mistakes to Avoid

1. Storing Calculated Values

sql

-- Don't store what you can calculate

CREATE TABLE Orders (

```
Quantity INT,  
UnitPrice DECIMAL(10,2),  
TotalPrice DECIMAL(10,2)    -- Will become inconsistent!  
);
```

-- Calculate on the fly or use computed column

CREATE TABLE Orders (

```
Quantity INT,  
UnitPrice DECIMAL(10,2),  
TotalPrice AS (Quantity * UnitPrice) PERSISTED  
);
```

2. Using VARCHAR for Numbers

sql

-- Wrong data type

```
CREATE TABLE Products (
    ProductID VARCHAR(50),      -- Should be INT
    Price VARCHAR(20),          -- Should be DECIMAL
    Quantity VARCHAR(10)        -- Should be INT
);
```

-- Correct data types

```
CREATE TABLE Products (
    ProductID INT,
    Price DECIMAL(10,2),
    Quantity INT
);
```

3. Not Considering NULL

sql

-- Allows NULL where it shouldn't

```
CREATE TABLE Users (
    UserID INT PRIMARY KEY,
    Email VARCHAR(100)      -- Email can be NULL?
);
```

-- Enforce NOT NULL

```
CREATE TABLE Users (
    UserID INT PRIMARY KEY,
    Email VARCHAR(100) NOT NULL -- Email is required
);
```

5.4 Performance Considerations

1. Index on Appropriate Columns

sql

-- Columns frequently used in WHERE, JOIN, ORDER BY

```
CREATE INDEX IX_Books_ISBN ON Books(ISBN);
CREATE INDEX IX_Loans_MemberID ON Loans(MemberID);
CREATE INDEX IX_Loans_DueDate ON Loans(DueDate);
```

2. Use Appropriate String Length

sql

```
-- Longer strings = slower indexing and comparisons  
-- Use realistic maximum lengths  
Email VARCHAR(100),          -- NOT VARCHAR(1000)  
Name VARCHAR(100),           -- NOT VARCHAR(500)  
PhoneNumber VARCHAR(15)       -- NOT VARCHAR(50)
```

3. Consider Storage Size

sql

```
-- Each row size affects performance  
-- 1 million rows with BIGINT = 8MB just for one column  
-- 1 million rows with INT = 4MB for same column  
-- Choose appropriately!
```

Quick Reference Summary

Data Type Cheat Sheet

For Text:

- Short text (< 50 chars): `VARCHAR(n)`
- Long text (variable): `VARCHAR(MAX)` or `NVARCHAR(MAX)`
- Fixed length (codes): `CHAR(n)`
- International text: `NVARCHAR(n)`

For Numbers:

- Small counts (0-255): `TINYINT`
- IDs, quantities: `INT`
- Very large numbers: `BIGINT`
- **Money (always!):** `DECIMAL(10,2)` or `DECIMAL(18,2)`
- Measurements: `DECIMAL(p,s)`
- Scientific: `FLOAT`
- Yes/No flags: `BIT`

For Dates/Times:

- Date only: `DATE`
- Time only: `TIME`
- Timestamp: `DATETIME2` (preferred) or `DATETIME`
- With timezone: `DATETIMEOFFSET`

For Special Cases:

- Unique IDs (distributed): `UNIQUEIDENTIFIER`
 - Files/images: `VARBINARY(MAX)`
 - XML data: `XML`
 - JSON data: `NVARCHAR(MAX)`
-

💡 Key Takeaways

1. **Choose data types carefully** - they affect storage, performance, and data integrity
 2. **Always use DECIMAL for money** - never FLOAT or REAL
 3. **Use smallest appropriate type** - saves space and improves performance
 4. **Consider international support** - use NVARCHAR for multilingual text
 5. **Use DATETIME2 over DATETIME** - better range and precision
 6. **Apply NOT NULL constraints** - prevent invalid data
 7. **Normalize properly** - aim for 3NF in most cases
 8. **Understand relationships** - use appropriate foreign keys
 9. **Index strategically** - improve query performance
 10. **Think about future needs** - but don't over-engineer
-

🎯 Practice Exercises

1. **Design a database** for an e-commerce system with appropriate data types
2. **Normalize** a denormalized table to 3NF

3. **Compare storage** requirements of different data types
 4. **Convert legacy DATETIME** columns to DATETIME2
 5. **Create indexes** for commonly queried columns
 6. **Implement constraints** to ensure data integrity
 7. **Design a multi-currency** pricing system using DECIMAL
 8. **Handle time zones** in a global application using DATETIMEOFFSET
-

Additional Resources

Official Documentation

- **Microsoft SQL Server Docs:** docs.microsoft.com/sql
- **MySQL Documentation:** dev.mysql.com/doc
- **PostgreSQL Documentation:** postgresql.org/docs
- **Oracle Database Docs:** docs.oracle.com/database

Learning Platforms

- **SQLZoo:** Interactive SQL tutorials
- **LeetCode Database:** Practice SQL problems
- **HackerRank SQL:** Coding challenges
- **Mode Analytics SQL Tutorial:** Real-world examples

Books

- "SQL in 10 Minutes, Sams Teach Yourself" by Ben Forta
- "Learning SQL" by Alan Beaulieu
- "Database Design for Mere Mortals" by Michael J. Hernandez
- "SQL Performance Explained" by Markus Winand

Video Courses

- **Microsoft Learn:** Free SQL courses
- **Coursera:** Database management courses

- **Udemy:** Comprehensive SQL courses
 - **YouTube:** FreeCodeCamp, Programming with Mosh
-

Glossary

ACID: Atomicity, Consistency, Isolation, Durability - properties that ensure database transaction reliability

Anomaly: Data inconsistency problem caused by poor database design (insert, update, delete anomalies)

Atomic Value: A value that cannot be divided into smaller parts (required for 1NF)

Candidate Key: A column or set of columns that could serve as the primary key

Composite Key: A primary key made up of two or more columns

Constraint: A rule enforced on data columns to maintain data integrity

DDL: Data Definition Language - SQL commands that define database structure (CREATE, ALTER, DROP)

DML: Data Manipulation Language - SQL commands that manipulate data (SELECT, INSERT, UPDATE, DELETE)

Foreign Key: A column that references the primary key of another table

Index: A database object that improves query performance

Junction Table: A table that connects two tables in a many-to-many relationship

Normalization: Process of organizing data to reduce redundancy and improve integrity

NULL: Absence of a value (different from zero or empty string)

Primary Key: A column or set of columns that uniquely identifies each row in a table

RDBMS: Relational Database Management System

Referential Integrity: Ensures that relationships between tables remain consistent

Schema: The structure of a database (tables, columns, relationships, constraints)

SQL: Structured Query Language - standard language for managing relational databases

Transitive Dependency: When a non-key column depends on another non-key column

Tuple: A row in a relational database table

Certification Paths

Entry Level

- **Microsoft:** Microsoft Certified: Azure Database Administrator Associate
- **Oracle:** Oracle Database SQL Certified Associate
- **MySQL:** MySQL 8.0 Database Administrator

Advanced Level

- **Microsoft:** Microsoft Certified: Azure Database Administrator Expert
 - **Oracle:** Oracle Database Administrator Certified Professional
 - **PostgreSQL:** PostgreSQL Certified Professional
-

Next Steps in Your SQL Journey

Beginner → Intermediate

1. Master basic SELECT queries
2. Learn JOINs (INNER, LEFT, RIGHT, FULL)
3. Understand GROUP BY and aggregate functions
4. Practice WHERE clause conditions
5. Learn subqueries and CTEs

Intermediate → Advanced

1. Study query optimization and indexing
2. Learn stored procedures and functions
3. Master transactions and concurrency
4. Understand execution plans
5. Practice complex queries with multiple JOINs

Advanced → Expert

1. Database administration and tuning

2. Replication and high availability
 3. Backup and recovery strategies
 4. Security and access control
 5. Data warehousing and ETL
-

Happy Learning! 

This document is designed to be a comprehensive reference for database management and SQL. Bookmark it, refer to it often, and practice the concepts with real projects. Remember: The best way to learn SQL is by doing!

Document Version: 1.0

Last Updated: December 2024

Created By: Database Training Team

Purpose: Comprehensive DBMS & SQL Training Material

© 2024 - This document is for educational purposes. Feel free to share and use for learning!